

# Predicting the car price based on its production year and mileage

## 1. Problem formulation:

The problem we focused on is predicting the car price based on its characteristics. The dataset contains a large number of cars both new and used that were listed for sale in the otomoto.pl portal. We focused on one specific vehicle type - Audi A3 with the engine size of 2000cm<sup>3</sup>.

We chose this problem, because we are interested in purchasing a car in near future and the analysis of the data can help us rate if the car price of specific parameters is reasonable or not. Another use case is to apply this model to vehicles with different brands and characteristics and check how common the model is/

The chosen dataset is called "Poland cars for sale dataset (200k adverts)" and can be found under this link <https://www.kaggle.com/datasets/bartoszipeniak/poland-cars-for-sale-dataset>. This dataset was created by webscraping over 200,000 car offers from one of the largest car advertisement sites in Poland (otomoto). It contains 25 parameters listed below:

ID - unique ID of offer Price - value of the price Currency - currency of the price (mostly Polish złoty, but also some euro) Condition - new or used Vehicle\_brand - brand of vehicle in offer Vehicle\_model - model of vehicle in offer Vehicle\_generation - generation of vehicle in offer Vehicle\_version - version of vehicle in offer Production\_year - year of car production Mileage\_km - total distance that the car has driven in kilometers Power\_HP - car engine power in horsepower Displacement\_cm3 - car engine size in cubic centimeters Fuel\_type - car fuel type CO2\_emissions - car CO2 emissions in g/km Drive - type of car drive Transmission - type of car transmission Type - car body style Doors\_number - number of car doors Colour - car body color Origin\_country - country of origin of the car First\_owner - whether the owner is the first owner First\_registration\_date - date of first registration Offer\_publication\_date - date of publication of the offer Offer\_location - address provided by the issuer Features - listed car features (ABS, airbag, parking sensors e.t.c)

### DAG Diagram

Based on the data, we created a DAG diagram to describe what parameters affect the price and each other. We divided the data into categories - brand specification, usage specification, car characteristics and external appearance. We also draw the relation between CO2 emission and parameters such as displacement, fuel type and horse type, which affect both the emission and the price.

## 2. Data preprocessing

Imports:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import arviz as az
import seaborn as sns
import cmdstanpy
import pandas as pd
import numpy as np
from scipy import stats
from scipy.optimize import curve_fit
import matplotlib.pyplot as plt
import seaborn as sns
import os
from sklearn.preprocessing import MinMaxScaler
from fitter import Fitter, get_common_distributions, get_distributions
```

BINS = 20

```
/usr/local/lib/python3.9/site-packages/tqdm/auto.py:22: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
```

Functions

```
def price_plot(df, column_name, plot_trend = False):
    price = df["Price"]
    data = df[column_name]
    plt.figure()
    plt.plot(data, price, 'o')
    plt.xlabel(column_name)
    plt.ylabel("Price_PLN")
    if plot_trend:
        z = np.polyfit(data, price, 1)
        p = np.poly1d(z)
        print(f"Polyfit equation: {p}")
        plt.plot(data, p(data))
        plt.axvline(data.mean(), color="red")
        plt.axhline(price.mean(), color="red")
    plt.show()
```

Loading the data:

```
df = pd.read_csv("data/Car_sale_ads.csv")
list(df.columns)
df.head()
```

	Index	Price	Currency	Condition	Vehicle_brand	Vehicle_model	\
0	0	86200	PLN	New	Abarth	595	
1	1	43500	PLN	Used	Abarth	Other	
2	2	44900	PLN	Used	Abarth	500	
3	3	39900	PLN	Used	Abarth	500	
4	4	97900	PLN	New	Abarth	595	

	Vehicle_version	Vehicle_generation	Production_year	Mileage_km	...
0	NaN	NaN	2021	1.0	...
1	NaN	NaN	1974	59000.0	...
2	NaN	NaN	2018	52000.0	...
3	NaN	NaN	2012	29000.0	...
4	NaN	NaN	2021	600.0	...

	Transmission	Type	Doors_number	Colour	Origin_country
0	Manual	small_cars	3.0	gray	NaN
1	Manual	coupe	2.0	silver	NaN
2	Automatic	small_cars	3.0	silver	NaN
3	Manual	small_cars	3.0	gray	NaN
4	Manual	small_cars	3.0	blue	NaN

	First_registration_date	Offer_publication_date	\
0	NaN	04/05/2021	
1	NaN	03/05/2021	
2	NaN	03/05/2021	
3	NaN	30/04/2021	
4	NaN	30/04/2021	

	Offer_location	\
0	ul. Jubilerska 6 - 04-190 Warszawa, Mazowiecki...	
1	kanonierska12 - 04-425 Warszawa, Rembertów (Po...	
2	Warszawa, Mazowieckie, Białołęka	
3	Jaworzno, Śląskie	
4	ul. Gorzysława 9 - 61-057 Poznań, Nowe Miasto ...	

```

                                Features
0                                []
1                                []
2  ['ABS', 'Electric front windows', 'Drivers air...
3  ['ABS', 'Electric front windows', 'Drivers air...
4  ['ABS', 'Electrically adjustable mirrors', 'Pa...

[5 rows x 25 columns]

```

Unification of the price currency and selection of the desired columns

```

price = df["Price"].copy()
currency = df["Currency"].copy()

for idx, (p, c) in enumerate(zip(price, currency)):
    if c == "EUR":
        price_PLN = p * 4.6
        price[idx] = price_PLN
        currency[idx] = "PLN"

df["Currency"] = currency
df["Price"] = price

cols2add = ["Price", "Vehicle_brand", "Vehicle_model",
            "Production_year", "Mileage_km", "Power_HP", "Displacement_cm3"]
test_df = df[cols2add]
test_df.head()

```


	Price	Vehicle_brand	Vehicle_model	Production_year	Mileage_km
0	86200.0	Abarth	595	2021	1.0
1	43500.0	Abarth	Other	1974	59000.0
2	44900.0	Abarth	500	2018	52000.0
3	39900.0	Abarth	500	2012	29000.0
4	97900.0	Abarth	595	2021	600.0


	Displacement_cm3
0	1400.0
1	1100.0
2	1368.0
3	1368.0
4	1368.0

Due to the extensive size of the dataset and the wide range of car models included, we have made the decision to conduct our analysis solely on a single car model. **Chosen car model:**  
**Brand:** Audi **Model:** A3


```
audi_cars = test_df[test_df['Vehicle_brand'] == "Audi"]  
audi_a3_cars = audi_cars[audi_cars["Vehicle_model"] == 'A3']  
price_plot(audi_a3_cars, "Displacement_cm3")
```



```
price_plot(audi_a3_cars, "Power_HP")
```



```
correlations = audi_a3_cars.iloc[:,  
1:].corrwith(audi_a3_cars['Price'])  
print(correlations)  
  
Production_year    0.853472  
Mileage_km         -0.764658  
Power_HP           0.354174  
Displacement_cm3   -0.301169  
dtype: float64  
  
correlation_matrix = audi_a3_cars.corr()  
mask = np.triu(np.ones_like(correlation_matrix, dtype=bool))  
sns.heatmap(data=correlation_matrix, mask=mask, annot=True,  
cmap='RdYlBu')  
plt.title('Lower Triangular Correlation Matrix')  
plt.show()
```



Due to small effect of engine power and displacement on the price of a vehicle, it was decided that only cars with a displacement of 2000ccm would be analysed to simplify analysis.

```
audi_a3_2010 = audi_a3_cars[audi_a3_cars["Production_year"] == 2010]
audi_a3_2000ccm = audi_a3_cars[audi_a3_cars["Displacement_cm3"] >= 1950]
audi_a3_2000ccm = audi_a3_2000ccm[audi_a3_2000ccm["Displacement_cm3"] <= 2050]
audi_a3_2000ccm = audi_a3_2000ccm.dropna()

if "audi_cars_data.csv" not in os.listdir("data"):
    audi_a3_2000ccm.to_csv('data/audi_cars_data.csv', index=False)
```

## Summary

```
audi_a3_2000ccm.head()
```

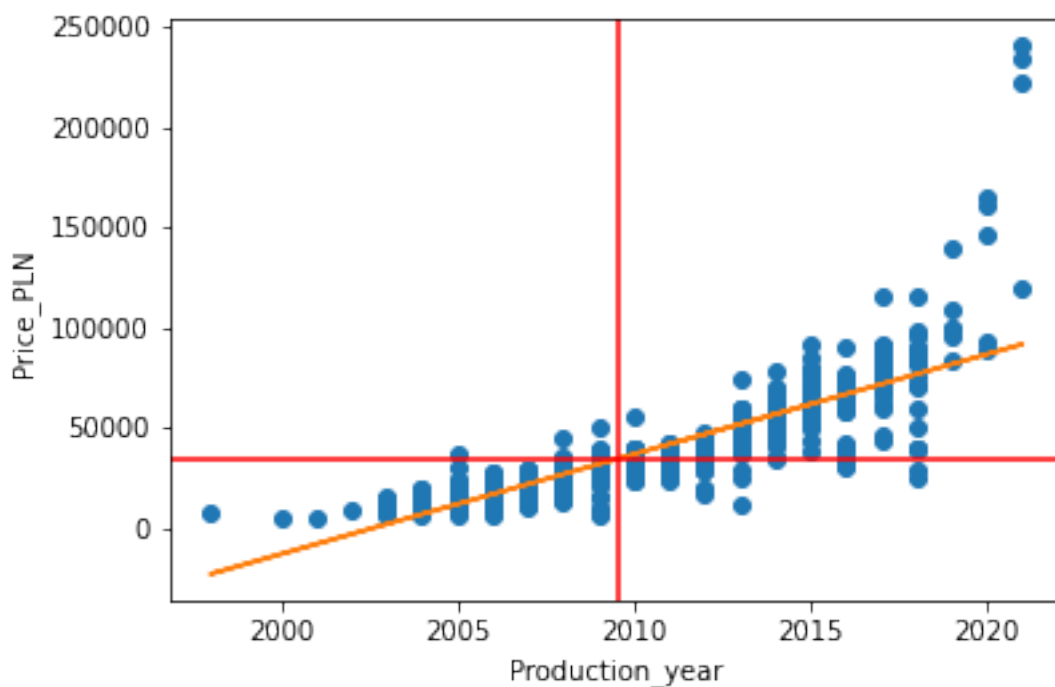
	Price	Vehicle_brand	Vehicle_model	Production_year	Mileage_km
1929	49900.0	Audi	A3	2015	208000.0
1932	13900.0	Audi	A3	2008	227000.0

1933	21900.0	Audi	A3	2008	313855.0
1934	19900.0	Audi	A3	2007	242000.0
1936	22900.0	Audi	A3	2006	240000.0

	Power_HP	Displacement_cm3
1929	150.0	1968.0
1932	140.0	1968.0
1933	140.0	1968.0
1934	170.0	1968.0
1936	200.0	1984.0

```
price_plot(audi_a3_2000ccm, "Production_year", True)
```


Polyfit equation:  
 $4988 x - 9.99e+06$



```
price_plot(audi_a3_2000ccm, "Mileage_km", True)
```


Polyfit equation:  
 $-0.2344 x + 8.109e+04$





```
price_plot(audi_a3_2000ccm, "Power_HP", True)
```

Polyfit equation:  
 $487.6 x - 4.072e+04$

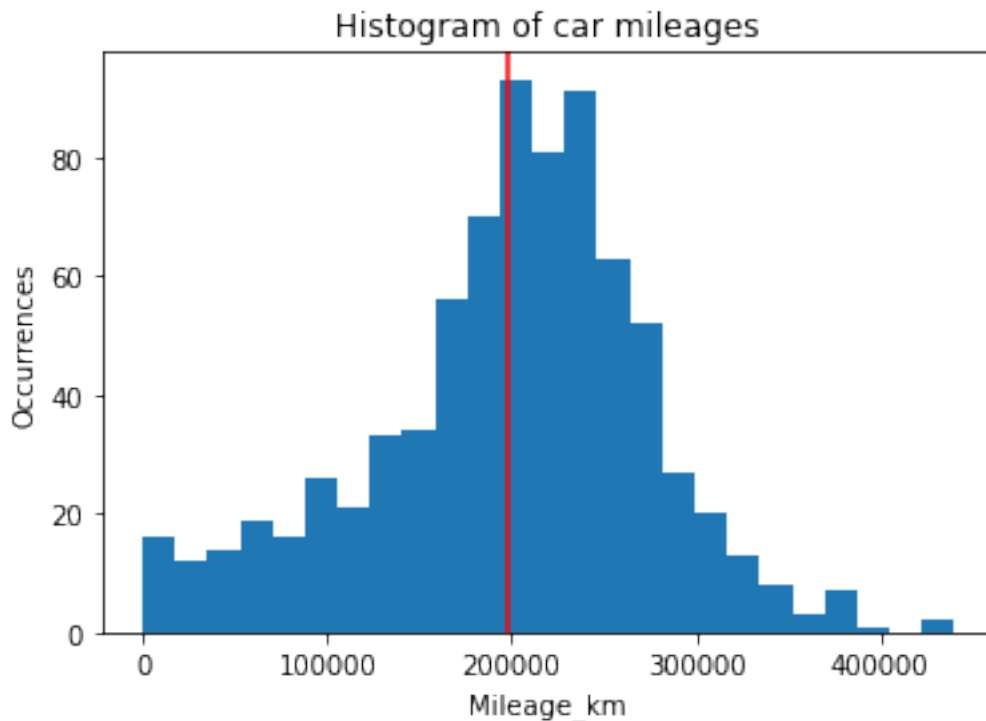


```

mileage_mean = np.mean(audi_a3_2000ccm["Mileage_km"])
print(f"Mean: {mileage_mean}")
plt.figure()
plt.hist(audi_a3_2000ccm["Mileage_km"], bins = 25)
plt.axvline(mileage_mean, color="red")
plt.xlabel("Mileage_km")
plt.ylabel("Occurrences")
plt.title("Histogram of car mileages")
plt.show()

```

Mean: 198361.37403598972




```

prod_mean = np.mean(audi_a3_2000ccm["Production_year"])
print(f"Mean: {prod_mean}")
plt.figure()
plt.hist(audi_a3_2000ccm["Production_year"], bins = 18)
plt.axvline(prod_mean, color="red")
plt.xlabel("Production year")
plt.ylabel("Occurrences")
plt.title("Histogram of car production years")
plt.show()


```

Mean: 2009.5719794344473



```
price_mean = np.mean(audi_a3_2000ccm["Price"])
print(f"Mean: {price_mean}")
plt.figure()
plt.hist(audi_a3_2000ccm["Price"], bins = 25)
plt.axvline(price_mean, color="red")
plt.xlabel("Price")
plt.ylabel("Occurrences")
plt.title("Histogram of cars price")
plt.show()
```

Mean: 34600.235218509



## Data standarization


Due to the diversity of the data (production year (values form 2003 to 2021), mileage (values from 0 to 400000), price (values from 0 to 160000)), we decided to standardise the data using the MinMax scaler. This way we got all the data in the range from 0 to 1, without losing information about data and making it easier to analyze it.

```
scaler = MinMaxScaler()
audi_a3_2000ccm_standarized_data =
scaler.fit_transform(audi_a3_2000ccm.loc[:,["Price",
"Production_year", "Mileage_km"]])
audi_a3_2000ccm_standarized =
pd.DataFrame(audi_a3_2000ccm_standarized_data,columns=["Price",
"Production_year", "Mileage_km"])
audi_a3_2000ccm_standarized.describe()
```


	Price	Production_year	Mileage_km
count	778.000000	778.000000	778.000000
mean	0.125764	0.503130	0.450820
std	0.110705	0.190895	0.173469
min	0.000000	0.000000	0.000000
25%	0.053109	0.347826	0.363635
50%	0.093048	0.478261	0.470453
75%	0.160497	0.652174	0.561363
max	1.000000	1.000000	1.000000

```
price_plot(audi_a3_2000ccm_standardized,"Production_year",True)  
price_plot(audi_a3_2000ccm_standardized,"Mileage_km",True)
```

Polyfit equation:  
 $0.4875 x - 0.1195$



Polyfit equation:  
 $-0.4382 x + 0.3233$



```

mileage_mean = np.mean(audi_a3_2000ccm_standardized["Mileage_km"])
print(f"Mean: {mileage_mean}")
plt.figure()
plt.hist(audi_a3_2000ccm_standardized["Mileage_km"], bins = 25)
plt.axvline(mileage_mean, color="red")
plt.xlabel("Mileage_km")
plt.ylabel("Occurrences")
plt.title("Histogram of car mileages")
plt.show()


prod_mean = np.mean(audi_a3_2000ccm_standardized["Production_year"])
print(f"Mean: {prod_mean}")
plt.figure()
plt.hist(audi_a3_2000ccm_standardized["Production_year"], bins = 18)
plt.axvline(prod_mean, color="red")
plt.xlabel("Production year")
plt.ylabel("Occurrences")
plt.title("Histogram of car production years")
plt.show()

price_mean = np.mean(audi_a3_2000ccm_standardized["Price"])
price_var = np.var(audi_a3_2000ccm_standardized["Price"])
print(f"Mean: {price_mean}")
print(f"Var: {price_var}")
plt.figure()
plt.hist(audi_a3_2000ccm_standardized["Price"], bins = 25)
plt.axvline(price_mean, color="red")
plt.xlabel("Price")

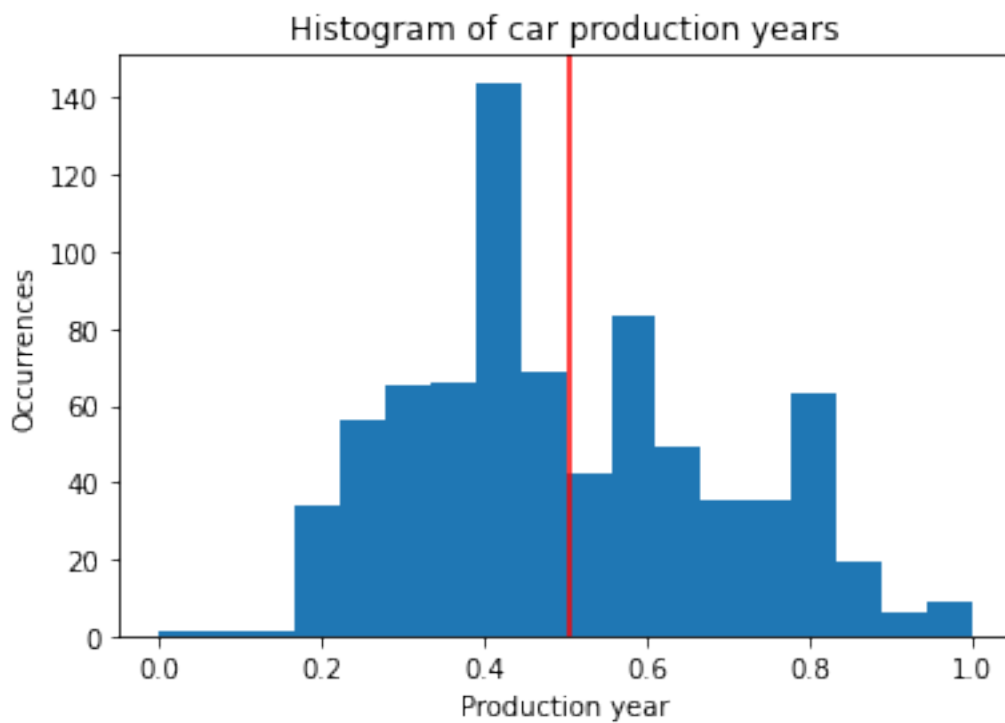
```

```
plt.ylabel("Occurrences")  
plt.title("Histogram of cars price")  
plt.show()
```

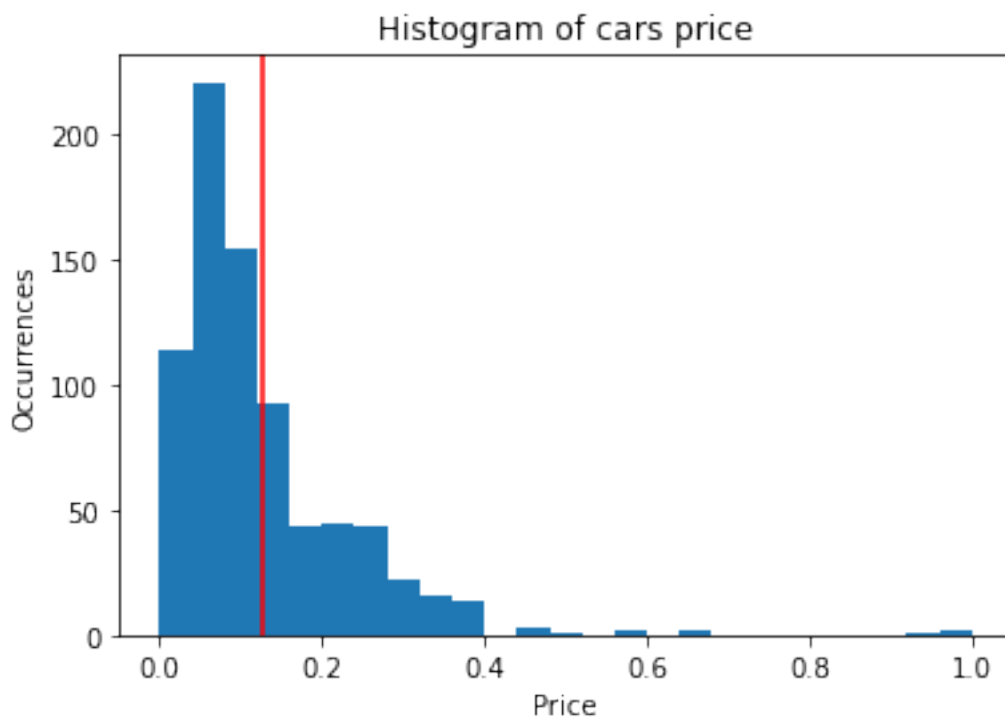
Mean: 0.45082005649101414



Mean: 0.5031295406281482



Mean: 0.12576418221431998  
Var: 0.012239734156792626





```

if 'audi_data_standardized.csv' not in os.listdir("data"):
    audi_a3_2000ccm_standardized.to_csv('data/audi_data_standardized.csv',
    index=False)

audi_a3_2000ccm_standardized =
pd.read_csv("data/audi_data_standardized.csv")
audi_a3_2000ccm_standardized.head()

```

	Price	Production_year	Mileage_km
0	0.190769	0.739130	0.472726
1	0.037814	0.434783	0.515908
2	0.071804	0.434783	0.713306
3	0.063306	0.391304	0.549999
4	0.076053	0.347826	0.545454

### 3. Model

For this project we specified two prior models of exponential range distribution. We wanted to check how adding highly correlated parameter to the model will affect price estimation. In the first model we used linear regression model with exponential distribution. We estimate price only based on the production year. In the second model we add mileage as well. Model 1 formula:

$$price = \text{exponential}((\alpha + \beta * \text{production\_year}) * \lambda)$$

Model 2 formula:

$$price = \text{exponential}((\alpha - \beta_1 * \text{mileage} + \beta_2 * \text{production\_year}) * \lambda)$$

#### 3.1 Model 1- prior

**Priors selection** The choice of an exponential distribution for modeling used car prices is justified by the observation that newer and less used cars tend to have higher prices. The exponential distribution captures this pattern with its right-skewed shape, accommodating a higher concentration of lower-priced cars and a smaller number of higher-priced cars. The min-max scaling ensures that the production year variable is on a comparable scale for accurate analysis and modeling of the relationship between production year and used car prices.

The choice to use a normal distribution for  $\alpha$ ,  $\beta$  and  $\lambda$  allows for capturing the natural variability of these parameters and is a common approach in statistical modeling for estimation and significance assessment.

```

model_exp1_ppc =
cmdstanpy.CmdStanModel(stan_file='stan_files/exp_model1_ppc.stan')

```

```

INFO:cmdstanpy:compiling stan file
/home/DA/project/stan_files/exp_model1_ppc.stan to exe file

```

```
/home/DA/project/stan_files/exp_model1_ppc
INFO:cmdstanpy:compiled model executable:
/home/DA/project/stan_files/exp_model1_ppc
```

```
N = len(audi_a3_2000ccm_standardized)
```

```
data = {"N": N,
        "mileage" : np.linspace(0.01,1,N),
        "production_year" : np.linspace(0.01,1,N)
}
```

```
sim_exp_fit1=model_exp1_ppc.sample(data=data)
```

```
sim_exp_fit1_pd = sim_exp_fit1.draws_pd()
```

```
sim_exp_fit1_pd.head()
```

```
INFO:cmdstanpy:CmdStan start processing
```

```
chain 1 | | 00:00 Status
```

```
| | 00:00 Status
```



```
| 00:00 Iteration: 200 / 1000 [ 20%] (Sampling)
```

```
| 00:00 Iteration: 400 / 1000 [ 40%] (Sampling)
```

```
| 00:00 Iteration: 600 / 1000 [ 60%] (Sampling)
```

```
| 00:00 Iteration: 800 / 1000 [ 80%] (Sampling)
```

```
chain 1 | | 00:00 Sampling completed
```

```
chain 2 | | 00:00 Sampling completed
```

```
chain 3 | | 00:00 Sampling completed
```

```
chain 4 | | 00:00 Sampling completed
```

```
INFO:cmdstanpy:CmdStan done processing.
```

lp__	accept_stat__	price[1]	price[2]	price[3]	price[4]
price[5]	\				
0	0.0	0.0	0.720541	0.056796	0.107119
0.011396					
1	0.0	0.0	0.103551	0.145529	0.402376
0.215753					
2	0.0	0.0	0.365362	0.128006	0.125368
0.142405					
3	0.0	0.0	0.057789	0.087359	0.031872
0.461796					
4	0.0	0.0	0.227178	0.047124	0.100616
0.001054					

price[6]	price[7]	price[8]	...	price[773]	price[774]
price[775]	\				

0	0.129340	0.213175	0.204050	...	0.045107	0.066630
	0.000519					
1	0.045607	0.066948	0.142049	...	0.000895	0.161457
	0.068008					
2	0.145356	0.041919	0.051312	...	0.091015	0.076533
	0.053217					
3	0.073865	0.032934	0.418010	...	0.038514	0.015578
	0.055715					
4	0.033103	0.116288	0.151095	...	0.030023	0.027987
	0.101090					

	price[776]	price[777]	price[778]	alpha	beta	sigma
lambda						
0	0.009614	0.168518	0.076745	0.189841	0.386983	0.154992
	40.0332					
1	0.205782	0.012172	0.000081	0.164562	0.341453	0.154589
	40.1370					
2	0.016594	0.052487	0.047103	0.167710	0.355258	0.132483
	39.9760					
3	0.007144	0.040562	0.003526	0.153393	0.370996	0.178134
	40.1906					
4	0.049186	0.004391	0.003777	0.199538	0.354305	0.189200
	39.6609					

[5 rows x 784 columns]


```
_, ax = plt.subplots(1, 4, figsize=(24, 5))
ax = ax.flatten()
sns.histplot(data=sim_exp_fit1_pd, x="alpha", stat="density",
ax=ax[0], bins=BINS)
sns.histplot(data=sim_exp_fit1_pd, x="beta", stat="density", ax=ax[1],
bins=BINS)
sns.histplot(data=sim_exp_fit1_pd, x="price[1]", stat="density",
ax=ax[2], bins=BINS)

ax[3].hist(sim_exp_fit1_pd["price[1]"], bins=BINS, alpha=0.5,
density=True, label="Prior")
ax[3].hist(audi_a3_2000ccm_standarized["Price"], bins=BINS, alpha=0.5,
density=True, label="Model samples")

ax[0].grid()
ax[1].grid()
ax[2].grid()
ax[3].grid()

ax[0].set_xlabel("alpha"),
ax[1].set_xlabel("beta"),
ax[2].set_xlabel("Predicted prices"),
ax[3].set_xlabel("Predicted prices vs dataset prices")
```

```
ax[3].set_ylabel("Density")
ax[3].legend()
plt.suptitle("Histograms")
plt.show()
```



The prior parameters were chosen through a semi-empirical process. Initially, a standard parameter from the literature was used, but further modifications were made to align the simulated data with the observed data. This adjustment ensured a closer match between the chosen priors and the actual data.

## 3.2 Model 1- posterior

```
model_exp_fit =
cmdstanpy.CmdStanModel(stan_file='stan_files/exp_model1_fit.stan')
N = len(audi_a3_2000ccm_standardized)
#Parameters
```

```
data = {"N": N,
        "mileage" : audi_a3_2000ccm_standardized['Mileage_km'],
        "production_year" :
audi_a3_2000ccm_standardized['Production_year'],
        "price_observed": audi_a3_2000ccm_standardized['Price']
}
```

```
sim_exp_pos1_fit=model_exp_fit.sample(data=data)
sim_exp_pos1_fit_pd = sim_exp_pos1_fit.draws_pd()
sim_exp_pos1_fit_pd.head()
```

```
INFO:cmdstanpy:compiling stan file
/home/DA/project/stan_files/exp_model1_fit.stan to exe file
/home/DA/project/stan_files/exp_model1_fit
INFO:cmdstanpy:compiled model executable:
/home/DA/project/stan_files/exp_model1_fit
INFO:cmdstanpy:CmdStan start processing
chain 1 | | 00:00 Status
| | 00:00 Status
| | 00:00 Iteration: 1 / 2000 [ 0%] (Warmup)
| | 00:02 Iteration: 100 / 2000 [ 5%] (Warmup)
```

```

chain 1 | ████████ | 00:02 Iteration: 300 / 2000 [ 15%] (Warmup)
chain 1 | ████████ | 00:02 Iteration: 400 / 2000 [ 20%] (Warmup)
chain 1 | ████████ | 00:03 Iteration: 600 / 2000 [ 30%] (Warmup)

chain 1 | ████████ | 00:03 Iteration: 800 / 2000 [ 40%] (Warmup)
chain 1 | ████████ | 00:03 Iteration: 900 / 2000 [ 45%] (Warmup)

██████████ | 00:04 Iteration: 1001 / 2000 [ 50%] (Sampling)
██████████ | 00:04 Iteration: 1100 / 2000 [ 55%] (Sampling)
██████████ | 00:05 Iteration: 1200 / 2000 [ 60%] (Sampling)
██████████ | 00:05 Iteration: 1300 / 2000 [ 65%] (Sampling)

██████████ | 00:06 Sampling completed
chain 2 | ██████████ | 00:06 Sampling completed
chain 3 | ██████████ | 00:06 Sampling completed
chain 4 | ██████████ | 00:06 Sampling completed

```

INFO:cmdstanpy:CmdStan done processing.

	lp__	accept_stat__	stepsize__	treedepth__	n_leapfrog__
divergent__					
0	692.464	0.995169	0.508152	3.0	7.0
0.0					
1	692.602	0.792205	0.508152	2.0	7.0
0.0					
2	690.839	0.643953	0.508152	2.0	3.0
0.0					
3	691.083	0.983516	0.508152	3.0	7.0
0.0					
4	692.512	0.976215	0.508152	2.0	7.0
0.0					

	energy__	alpha	beta	sigma	...	log_lik[769]
log_lik[770]						
0	-690.924	0.149932	0.104850	0.133294	...	1.12887
1.48035						
1	-690.797	0.151604	0.100408	0.146773	...	1.12891
1.48533						
2	-690.241	0.159595	0.109254	0.153538	...	1.12729
1.51199						
3	-690.081	0.145114	0.088603	0.143667	...	1.12615
1.45404						

```
4 -690.673 0.148307 0.116666 0.138012 ... 1.12874
1.48937
```

	log_lik[771]	log_lik[772]	log_lik[773]	log_lik[774]
0	1.13240	1.65627	0.889847	1.26490
1	1.13247	1.66221	0.888445	1.26490
2	1.13107	1.69850	0.873103	1.26303
3	1.12946	1.62042	0.900512	1.26163
4	1.13238	1.66967	0.883629	1.26423

	log_lik[776]	log_lik[777]	log_lik[778]
0	1.38526	1.45391	0.884999
1	1.38862	1.45404	0.883625
2	1.40826	1.46475	0.866488
3	1.36419	1.43652	0.897635
4	1.39279	1.46117	0.877504

```
[5 rows x 2345 columns]
```


```
_, ax = plt.subplots(1, 4, figsize=(24, 5))
ax = ax.flatten()
sns.histplot(data=sim_exp_pos1_fit_pd, x="alpha", stat="density",
ax=ax[0], bins=BINS)
sns.histplot(data=sim_exp_pos1_fit_pd, x="beta", stat="density",
ax=ax[1], bins=BINS)
sns.histplot(data=sim_exp_pos1_fit_pd, x="price_estimated[1]",
stat="density", ax=ax[2], bins=BINS)

ax[3].hist(sim_exp_pos1_fit_pd["price_estimated[1]"], bins=BINS,
alpha=0.5, density=True, label="Posterior")
ax[3].hist(audi_a3_2000ccm_standarized["Price"], bins=BINS, alpha=0.5,
density=True, label="Model samples")

ax[0].grid()
ax[1].grid()
ax[2].grid()
ax[3].grid()

ax[0].set_xlabel("alpha"),
ax[1].set_xlabel("beta"),
ax[2].set_xlabel("Predicted prices"),
ax[3].set_xlabel("Predicted prices and model samples comparison")
```

```
ax[3].set_ylabel("Density")
ax[3].legend()
plt.suptitle("Histograms")
plt.show()
```



The posterior data analysis, which includes comparing the histogram of the prior distribution with the posterior distribution and real data, indicates a good fit. This comparison demonstrates that the chosen priors align well with the observed data, suggesting that the model captures the underlying patterns and provides reliable estimates.

```
summary = sim_exp_pos1_fit.summary()
summary.head()
```

	Mean	MCSE	StdDev	5%	50%	95%	N_Eff
N_Eff/s \							
name							
lp__	690.00	0.03000	1.4000	690.000	690.00	690.00	2100.0
140.0							
alpha	0.15	0.00018	0.0099	0.140	0.15	0.17	3100.0
210.0							
beta	0.10	0.00031	0.0170	0.074	0.10	0.13	3000.0
200.0							
sigma	0.15	0.00033	0.0200	0.120	0.15	0.18	3500.0
240.0							
lambda	40.00	0.00340	0.2000	40.000	40.00	40.00	3500.0
240.0							
R_hat							
name							
lp__	1.0						
alpha	1.0						
beta	1.0						
sigma	1.0						
lambda	1.0						

### 3.3 Model 2- prior

The extension of the first model to include mileage introduces an additional predictor variable, expanding the model's scope. This extension allows for the consideration of mileage as a factor influencing used car prices. By incorporating mileage into the model, it is possible to assess its impact on the relationship between other predictors (such as production year) and used car prices.

In the extended model, the inclusion of mileage as a predictor involved adding a normal distribution parameter, "beta1". Other parameters stayed the same.

```
model_exp2_ppc =  
cmdstanpy.CmdStanModel(stan_file='stan_files/exp_model2_ppc.stan')
```

```
#Parameters
```

```
N = len(audi_a3_2000ccm_standardized)
```

```
mu_a = 0.17
```

```
sig_a = 0.02
```

```
mu_b1 = 0.36
```

```
sig_b1 = 0.02
```

```
mu_b2 = 0.36
```

```
sig_b2 = 0.02
```

```
data = {"N": N,  
        "mileage" : np.linspace(0.01,1,N),  
        "production_year" : np.linspace(0.01,1,N),  
        "mu_a" : mu_a,  
        "sig_a" : sig_a,  
        "mu_b1" : mu_b1,  
        "mu_b2" : mu_b2,  
        "sig_b1" : sig_b1,  
        "sig_b2" : sig_b2,  
        }
```

```
sim_exp_fit2=model_exp2_ppc.sample(data=data)
```

```
sim_exp_fit2_pd = sim_exp_fit2.draws_pd()
```

```
sim_exp_fit2_pd.head()
```

```
INFO:cmdstanpy:compiling stan file
```

```
/home/DA/project/stan_files/exp_model2_ppc.stan to exe file
```

```
/home/DA/project/stan_files/exp_model2_ppc
```

```
INFO:cmdstanpy:compiled model executable:
```

```
/home/DA/project/stan_files/exp_model2_ppc
```

```
INFO:cmdstanpy:CmdStan start processing
```

```
chain 1 | | 00:00 Status
```

```
| 00:00 Status
```

```
| 00:00 Iteration: 100 / 1000 [ 10%] (Sampling)
```

```
| 00:00 Iteration: 300 / 1000 [ 30%] (Sampling)
```



```
██████████ | 00:00 Iteration: 500 / 1000 [ 50%] (Sampling)
```

```
██████████ | 00:00 Iteration: 700 / 1000 [ 70%] (Sampling)
```

```
██████████ | 00:00 Sampling completed
```

```
chain 2 | ██████████ | 00:00 Sampling completed
```

```
chain 3 | ██████████ | 00:00 Sampling completed
```

```
chain 4 | ██████████ | 00:00 Sampling completed
```

INFO:cmdstanpy:CmdStan done processing.

	lp__	accept_stat__	price[1]	price[2]	price[3]	price[4]
price[5]	\					
0	0.0	0.0	0.191955	0.381461	0.139544	0.114504
0.038685						
1	0.0	0.0	0.160646	0.329924	0.133558	0.106174
0.123089						
2	0.0	0.0	0.033113	0.115929	0.494742	0.014967
0.198760						
3	0.0	0.0	0.044376	0.045677	0.178827	0.074792
0.085814						
4	0.0	0.0	0.065762	0.082820	0.276127	0.501971
0.015260						

	price[6]	price[7]	price[8]	...	price[774]	price[775]
price[776]	\					
0	0.130348	0.102170	0.169845	...	0.003976	0.147585
0.007609						
1	0.284623	0.008743	0.036188	...	0.206796	0.098887
0.063346						
2	0.026657	0.390721	0.037631	...	0.320360	0.442588
0.641653						
3	0.381912	0.249573	0.078680	...	0.396380	0.029455
0.157756						
4	0.029156	0.056822	0.009745	...	0.232063	0.018306
0.066841						

	price[777]	price[778]	alpha	beta1	beta2	sigma
lambda						
0	0.134670	0.043093	0.197707	0.397723	0.336356	0.132193
40.0124						
1	0.215654	0.029304	0.155578	0.337800	0.357244	0.124951
40.3343						
2	0.483538	0.106727	0.183627	0.390463	0.327905	0.151000
39.7154						
3	0.339884	0.298564	0.155037	0.362080	0.356258	0.145441

```
40.1420
4    0.076361    0.019694    0.169357    0.344133    0.375236    0.166650
40.0501
```

```
[5 rows x 785 columns]
```


```
_, ax = plt.subplots(1, 5, figsize=(24, 5))
ax = ax.flatten()
sns.histplot(data=sim_exp_fit2_pd, x="alpha", stat="density",
ax=ax[0], bins=BINS)
sns.histplot(data=sim_exp_fit2_pd, x="beta1", stat="density",
ax=ax[1], bins=BINS)
sns.histplot(data=sim_exp_fit2_pd, x="beta2", stat="density",
ax=ax[2], bins=BINS)
sns.histplot(data=sim_exp_fit2_pd, x="price[1]", stat="density",
ax=ax[3], bins=BINS)

ax[4].hist(sim_exp_fit2_pd["price[1]"], bins=BINS, alpha=0.5,
density=True, label="Prior")
ax[4].hist(audi_a3_2000ccm_standarized["Price"], bins=BINS, alpha=0.5,
density=True, label="Model samples")

ax[0].grid()
ax[1].grid()
ax[2].grid()
ax[3].grid()

ax[0].set_xlabel("alpha"),
ax[1].set_xlabel("beta1"),
ax[2].set_xlabel("beta2"),
ax[3].set_xlabel("Predicted prices"),
ax[4].set_xlabel("Predicted prices vs dataset prices")

ax[4].set_ylabel("Density")
ax[4].legend()
plt.suptitle("Histograms")
plt.show()
```



The comparison between the prior model and real data suggests a satisfactory fit, indicating that the chosen priors accurately capture the underlying patterns and characteristics of the observed data.

## 3.4 Model 2- posterior

```
model_exp2_fit =  
cmdstanpy.CmdStanModel(stan_file='stan_files/exp_model2_fit.stan')  
N = len(audi_a3_2000ccm_standardized)  
#Parameters
```

```
data = {"N": N,  
        "mileage" : audi_a3_2000ccm_standardized['Mileage_km'],  
        "production_year" :  
audi_a3_2000ccm_standardized['Production_year'],  
        "price_observed": audi_a3_2000ccm_standardized['Price']  
}
```

```
sim_exp_pos2_fit=model_exp2_fit.sample(data=data)  
sim_exp_pos2_fit_pd = sim_exp_pos2_fit.draws_pd()  
sim_exp_pos2_fit_pd.head()
```

```
INFO:cmdstanpy:compiling stan file  
/home/DA/project/stan_files/exp_model2_fit.stan to exe file  
/home/DA/project/stan_files/exp_model2_fit  
INFO:cmdstanpy:compiled model executable:  
/home/DA/project/stan_files/exp_model2_fit  
INFO:cmdstanpy:CmdStan start processing  
chain 1 | | 00:00 Status  
█ | 00:00 Iteration: 100 / 2000 [ 5%] (Warmup)  
█ | 00:05 Iteration: 300 / 2000 [ 15%] (Warmup)
```

```
█ | 00:10 Sampling completed  
chain 2 | █ | 00:10 Sampling completed  
chain 3 | █ | 00:10 Sampling completed  
chain 4 | █ | 00:10 Sampling completed
```

```
INFO:cmdstanpy:CmdStan done processing.
```

	lp__	accept_stat__	stepsize__	treedepth__	n_leapfrog__
divergent__	\				
0	558.268	0.991764	0.398757	3.0	7.0
0.0					
1	558.577	0.986237	0.398757	3.0	15.0
0.0					
2	560.092	0.990274	0.398757	3.0	15.0

```

0.0
3  560.113      0.940848    0.398757      3.0      7.0
0.0
4  561.055      0.935519    0.398757      3.0      7.0
0.0

```

```

    energy__    alpha    beta1    beta2    ...    log_lik[769]
log_lik[770] \
0  -553.932  0.217813  0.156296  0.097847  ...      1.12817
1.49560
1  -556.565  0.210059  0.134993  0.120802  ...      1.12433
1.51474
2  -556.329  0.227846  0.143052  0.085959  ...      1.12577
1.52220
3  -557.540  0.216448  0.135953  0.098864  ...      1.12724
1.50716
4  -557.922  0.210008  0.140887  0.124191  ...      1.12420
1.51315

```

```

    log_lik[771]  log_lik[772]  log_lik[773]  log_lik[774]
log_lik[775] \
0      1.12833      1.58581      0.878486      1.24921
0.147678
1      1.13243      1.63329      0.862541      1.23857
0.061606
2      1.13182      1.63531      0.865778      1.24764
0.111412
3      1.13112      1.61775      0.872966      1.24987
0.121647
4      1.13238      1.62775      0.862380      1.23614
0.058876

```

```

    log_lik[776]  log_lik[777]  log_lik[778]
0      1.30377      1.47543      0.871042
1      1.34413      1.47431      0.852473
2      1.34361      1.47539      0.858188
3      1.33312      1.47544      0.865295
4      1.33879      1.47372      0.851951

```

```
[5 rows x 2346 columns]
```

```

_, ax = plt.subplots(1, 5, figsize=(24, 5))
ax = ax.flatten()
sns.histplot(data=sim_exp_pos2_fit_pd, x="alpha", stat="density",
ax=ax[0], bins=BINS)
sns.histplot(data=sim_exp_pos2_fit_pd, x="beta1", stat="density",
ax=ax[1], bins=BINS)
sns.histplot(data=sim_exp_pos2_fit_pd, x="beta2", stat="density",
ax=ax[2], bins=BINS)
sns.histplot(data=sim_exp_pos2_fit_pd, x="price_estimated[1]",

```

```

stat="density", ax=ax[3], bins=BINS)


ax[4].hist(sim_exp_pos2_fit_pd["price_estimated[1]"], bins=BINS,
alpha=0.5, density=True, label="Posterior")
ax[4].hist(audi_a3_2000ccm_standarized["Price"], bins=BINS, alpha=0.5,
density=True, label="Model samples")

ax[0].grid()
ax[1].grid()
ax[2].grid()
ax[3].grid()

ax[0].set_xlabel("alpha"),
ax[1].set_xlabel("beta1"),
ax[2].set_xlabel("beta2"),
ax[3].set_xlabel("Predicted prices"),
ax[4].set_xlabel("Predicted prices vs dataset prices")

ax[4].set_ylabel("Density")
ax[4].legend()
plt.suptitle("Histograms")
plt.show()

```



The posterior distribution of the model, after fitting it to the real data, exhibits a good fit, indicating that the model effectively captures the patterns and characteristics present in the observed data.

```

summary = sim_exp_pos2_fit.summary()
summary.head(6)

```

	Mean	MCSE	StdDev	5%	50%	95%	N_Eff
N_Eff/s \							
name							
lp__	560.00	0.03800	1.600	560.000	560.00	560.00	1800.0
120.0							
alpha	0.21	0.00025	0.012	0.190	0.21	0.23	2200.0
150.0							
beta1	0.13	0.00034	0.017	0.100	0.13	0.16	2300.0


150.0							
beta2	0.11	0.00031	0.017	0.079	0.11	0.13	3000.0
190.0							
sigma	0.15	0.00033	0.020	0.120	0.15	0.18	3500.0
230.0							
lambda	40.00	0.00280	0.200	40.000	40.00	40.00	4900.0
320.0							

	R_hat
name	
lp__	1.0
alpha	1.0
beta1	1.0
beta2	1.0
sigma	1.0
lambda	1.0

## 4. Model comparison

```
compare_model_waic = az.compare(
    {
        "fit_1": az.from_cmdstanpy(sim_exp_pos1_fit),
        "fit_2": az.from_cmdstanpy(sim_exp_pos2_fit)
    },
    ic="waic",
)

ax = az.plot_compare(compare_model_waic)
ax.set_title(f"Comparison of models with waic criterion")
plt.show()
```



```
display(compare_model_waic)
```

	rank	waic	p_waic	d_waic	weight	se
fit_1	0	773.471232	1.651359	0.000000	1.000000e+00	28.535902
fit_2	1	706.216197	2.543949	67.255035	8.954970e-11	32.926318

	dse	warning	waic_scale
fit_1	0.000000	False	log
fit_2	5.455465	False	log

The results of the WAIC (Watanabe-Akaike information criterion) comparison are as follows:


- rank: The ranking of the models based on their WAIC values. In this case, fit\_1 is ranked at 0, indicating that fit\_1 is preferred model
- waic: The WAIC value for each model. The WAIC is a measure of the out-of-sample predictive accuracy of the model. In this case, fit\_1 has a higher WAIC value of 773.187452, while fit\_2 has a lower WAIC value of 705.928822.
- p\_waic: The estimated effective number of parameters based on the WAIC, used to compare the complexity of the models.
- d\_waic: The difference in WAIC values between the models. In this case, fit\_2 has a higher WAIC value by 67.25863 compared to fit\_1.
- weight: The weight of each model in the model comparison, representing the probability of each model being the best model among the compared models. In this case, fit\_1 has a weight of 1.0, indicating it is the preferred model over fit\_2, which has a weight of 0.0.
- se: The standard error of the WAIC estimate, providing a measure of uncertainty associated with the WAIC value.
- dse: The standard error of the difference in WAIC values, providing a measure of uncertainty associated with the difference in WAIC values between the models.
- warning: Indicates whether there are any warnings associated with the model comparison. In this case, there is no warning. waic\_scale: The scale of the WAIC values. In this case, the values are on a log scale.

Based on these results, fit\_2 is ranked higher with a lower WAIC value, indicating better model performance. Although the weight for fit\_1 is 1.0, suggesting it is the preferred model, this contradicts the lower WAIC value of fit\_2. Additionally the rank suggest to choose fit\_1. Therefore, the results of fit\_1 and fit\_2 overlap and further analysis is needed

```
compare_model_loo = az.compare(
{
"fit_1": az.from_cmdstanpy(sim_exp_pos1_fit),
"fit_2": az.from_cmdstanpy(sim_exp_pos2_fit)
},
ic="loo",
)

ax = az.plot_compare(compare_model_loo)
```

```
ax.set_title(f"Comparison of models with loo criterion")
plt.show()
```



```
display(compare_model_loo)
```

	rank	loo	p_loo	d_loo	weight	se
fit_1	0	773.472448	1.650144	0.000000	1.000000e+00	28.535567
fit_2	1	706.214781	2.545365	67.257666	3.261391e-12	32.926727

	dse	warning	loo_scale
fit_1	0.000000	False	log
fit_2	5.456355	False	log

The results of the LOO analysis comparison are as follows:

- rank: The ranking of the models based on their LOO values. In this case, fit\_1 is ranked at 0, indicating that fit\_1 is preferred.
- loo: The LOO value for each model. The LOO is a measure of the out-of-sample predictive accuracy of the model. In this case, fit\_1 has a higher LOO value of 773.186874, while fit\_2 has a lower LOO value of 705.931121.
- p\_loo: The estimated effective number of parameters based on the LOO, used to compare the complexity of the models.
- d\_loo: The difference in LOO values between the models. In this case, fit\_2 has a higher LOO value by 67.255753 compared to fit\_1.
- weight: The weight of each model in the model comparison, representing the probability of each model being the best model among the compared models. In this case, fit\_1 has a weight of 1.0, indicating it is the preferred model over fit\_2, which has a weight of 0.0.
- se: The standard error of the LOO estimate, providing a measure of uncertainty associated with the LOO value.
- dse: The standard error of the difference in LOO values, providing a measure of uncertainty associated with the difference in LOO values between the models.



- `warning`: Indicates whether there are any warnings associated with the model comparison. In this case, there is no warning.
- `loo_scale`: The scale of the LOO values. In this case, the values are on a log scale.

Based on these results, `fit_2` is ranked higher with a lower WAIC value, indicating better model performance. Although the weight for `fit_1` is 1.0, suggesting it is the preferred model, this contradicts the lower LOO value of `fit_2`. Additionally the rank suggest to choose `fit_1`. Therefore, the results of `fit_1` and `fit_2` overlap and further analysis is needed